

Experiment 1: Titanic Dataset

COGS118A - Final Project

by Justin Laughlin and Stuart Sonatina

```
In [1]: # import packages
        from datetime import datetime
        import numpy as np
        import matplotlib.pyplot as plt
        import scipy.io as sio
        import pandas as pd
        import seaborn as sns
        sns.set(style="whitegrid", palette="muted")

        # splitting, training, and testing
        import sklearn.model_selection as ms
        #from sklearn.model_selection import train_test_split
        #from sklearn.model_selection import cross_val_score

        # Principle component analysis
        from sklearn.decomposition import PCA

        # random forest classifier
        from sklearn.ensemble import RandomForestClassifier

        # Visualization
        import sklearn.metrics as skm

        # magic command to display plots inline
        %matplotlib inline
```

Features of "titanic.csv":

Variable	Definition	Key
survival	Survival	0 = No, 1 = Yes
pclass	Ticket class	1 = 1st, 2 = 2nd, 3 = 3rd
sex	Sex	0 = Male, 1 = Female
Age	in years	
sibsp	# of siblings / spouses aboard the Titanic	
parch	# of parents / children aboard the Titanic	
ticket	Ticket number	
fare	Passenger fare	
cabin	Cabin number	
embarked	Port of Embarkation	C = Cherbourg, Q = Queenstown, S = Southampton

Source: [Kaggle \(https://www.kaggle.com/c/titanic/data\)](https://www.kaggle.com/c/titanic/data)

```
In [2]: # Setting parameters
trainTestSplit = [0.4, 0.6, 0.8]
nfold = 10
ntrees = 1000
# (https://matplotlib.org/users/colormaps.html)
colorMap = 'GnBu'
```

```
In [3]: # Use pandas to import data into dataframe. PassengerID is redundant
df = pd.read_csv('titanic.csv')
df.drop('PassengerId', axis=1, inplace=True)
```

```

In [4]: # class label == survival
Y = df.as_matrix(columns=['Survived']).reshape([-1,1]).ravel()

# Delete Name, Ticket and Cabin features, along with class label and
# encoded features
Xdf = df.drop(['Name', 'Ticket', 'Cabin', 'Survived', 'Embarked'], ax
is=1)

# Convert NaN to mean age
Xdf.fillna(np.mean(Xdf['Age']), inplace=True)

# One-hot encode port embarked from
PortEncoded = pd.get_dummies(df['Embarked'])
Xdf['EmbCherb'] = PortEncoded['C']
Xdf['EmbQueen'] = PortEncoded['Q']
Xdf['EmbSouth'] = PortEncoded['S']

# replace male/female with 0/1 respectively
Xdf['Sex'].replace(['male', 'female'], [0., 1.], inplace=True)

num_samples, num_features = Xdf.shape

```

```

In [5]: # Look at first 6 passengers
Xdf[:6]

```

```

Out[5]:

```

	Pclass	Sex	Age	SibSp	Parch	Fare	EmbCherb	EmbQueen	EmbSouth
0	3	0.0	22.000000	1	0	7.2500	0	0	1
1	1	1.0	38.000000	1	0	71.2833	1	0	0
2	3	1.0	26.000000	0	0	7.9250	0	0	1
3	1	1.0	35.000000	1	0	53.1000	0	0	1
4	3	0.0	35.000000	0	0	8.0500	0	0	1
5	3	0.0	29.699118	0	0	8.4583	0	1	0


Data Exploration

```
In [6]: # Male/Female vs age
# wideform dataframe
age_vs_sex = pd.DataFrame.from_dict({'Sex':Xdf['Sex'], 'Age':Xdf['Age'], 'Survived':df['Survived']})

# Draw a categorical scatterplot to show each observation
fig, ax = plt.subplots()
fig.set_size_inches(11,8)
sns.set()
sns.set(font_scale=1.5)
sns.swarmplot(x="Sex", y="Age", hue="Survived", data=age_vs_sex)
plt.legend(prop={'size':20})

print("Score for 'Female survived': {:.1f}%".format((np.sum(df['Survived']==Xdf['Sex'])/num_samples)*100))
```

Score for 'Female survived': 78.7%



```
In [7]: # Plot correlation matrix
plt.figure(figsize=(11,8))
sns.heatmap(df.corr(),square=True,annot=True,cmap="RdBu")
```

Out[7]: <matplotlib.axes._subplots.AxesSubplot at 0x7fb6c34abe48>



Feature Engineering

```
In [8]: # Save a copy of Xdf before feature engineering so it can be compared
Xdf_nofe = Xdf.copy(deep=True)
```

```
In [9]: # If individual is under 18 and Parch>=1 individual is a child with p
        # arent(s) on board
Xdf["ChwPar"] = ((Xdf["Age"]<=18) & (Xdf["Parch"]>=1)) * Xdf["Parch"]
        # If individual is over 18 and Parch>=1 individual is a parent with c
        # hildren on board
Xdf["ParwCh"] = ((Xdf["Age"]>18) & (Xdf["Parch"]>=1)) * Xdf["Parch"]
        # Drop the feature "Parch" as it has been split into two different fe
        # atures
Xdf = Xdf.drop(["Parch"], axis=1)
```

```

In [10]: # Extract ChwPar and ParwCh from dataframes
ChwPar = pd.DataFrame.from_dict({'# Parents':Xdf['ChwPar'], 'Survived':df['Survived']})
ParwCh = pd.DataFrame.from_dict({'# Children':Xdf['ParwCh'], 'Survived':df['Survived']})


# Setting up plot environment
alph = 5 # significance level [%] for confidence intervals
cs = 0.15 # capsize for confidence interval tips
fig, (ax1,ax2) = plt.subplots(1,2,sharey=True)
fig.set_size_inches(11,8)
ax1.set_ylim(0,1)
sns.set()
sns.set(font_scale=1.5)

# Draw a barplot for ChwPar
sns.barplot(x="# Parents", y="Survived", data=ChwPar, ax=ax1, color = "turquoise", ci=100-alph, capsize=cs)

# Draw a barplot for ParwCh
sns.barplot(x="# Children", y="Survived", data=ParwCh, ax=ax2, color = "turquoise", ci=100-alph, capsize=cs)

```


Out[10]: <matplotlib.axes._subplots.AxesSubplot at 0x7fb6c1b89be0>



In [14]: Xdf.head()

Out[14]:

	Pclass	Sex	Age	SibSp	Fare	EmbCherb	EmbQueen	EmbSouth	ChwPar	Parw
0	3	0.0	22.0	1	7.2500	0	0	1	0	0
1	1	1.0	38.0	1	71.2833	1	0	0	0	0
2	3	1.0	26.0	0	7.9250	0	0	1	0	0
3	1	1.0	35.0	1	53.1000	0	0	1	0	0
4	3	0.0	35.0	0	8.0500	0	0	1	0	0



Training Classifier

In [15]: Xdf_nofe.head()

Out[15]:

	Pclass	Sex	Age	SibSp	Parch	Fare	EmbCherb	EmbQueen	EmbSouth
0	3	0.0	22.0	1	0	7.2500	0	0	1
1	1	1.0	38.0	1	0	71.2833	1	0	0
2	3	1.0	26.0	0	0	7.9250	0	0	1
3	1	1.0	35.0	1	0	53.1000	0	0	1
4	3	0.0	35.0	0	0	8.0500	0	0	1


```

In [16]: # Function which performs a grid search to estimate the best parameters and
# plots the heatmap of scores vs. parameters
def gridSearchAndPlot(estimator, Xtrain, Ytrain):

    # grid search
    clf = ms.GridSearchCV(estimator, parameters, cv=nfold)
    clf.fit(Xtrain, Ytrain)
    # print(clf.best_params_)
    cvScores = clf.cv_results_["mean_test_score"].reshape([a,b])

    # plot
    plt.figure()
    plt.pcolor(parameters['min_samples_split'], parameters['max_features'], cvScores, cmap=colorMap)
    plt.colorbar()
    plt.xlabel('Min Samples to Split')
    plt.ylabel('Max number of Features Considered')
    plt.xticks(parameters['min_samples_split'])
    plt.yticks(parameters['max_features'])
    plt.title('Cross Validation Scores vs Hyperparameters')

    # return classifier
    return clf

```

```

In [17]: # F-score provides a useful balance between precision and recall
# By definition:  $F = 2 * (precision * recall) / (precision + recall)$ 
def calcFscore(clf, Xtest, Ytest):
    # Calculate score
    Yscore = clf.predict(Xtest)

    precision, recall, thresholds = skm.precision_recall_curve(Ytest, Yscore, pos_label=1)
    skm.classification_report(Ytest, Yscore, target_names=['Didn\'t Survive', 'Survived'])
    F = 2*(precision[1]*recall[1])/(precision[1]+recall[1])

    # Print results
    # print("Precision: \t{}".format(precision))
    # print("Recall: \t{}".format(recall))
    # print("Thresholds: \t{}".format(thresholds))
    # print(F)

    # return
    return (precision, recall, F)

```

```
In [18]: # Initialize stuff for classifier

# dictionary of classifiers and scores (precision, recall, and F-score)
clf = {}
scores = {"precision": {}, "recall": {}, "F": {}}

# train a random forest classifier on our training set
rand_forest = RandomForestClassifier(n_estimators=ntrees)

# use GridSearchCV to tune the following hyperparameters:
# min_samples_split == minimum samples in node to perform a split
# max_features == max number of features used to perform a split
parameters = {'min_samples_split': [2, 3, 4, 5, 6, 7, 8], 'max_features':
[1, 2, 3, 4, 5, 6, 7]}
b = len(parameters['min_samples_split'])
a = len(parameters['max_features'])
```

Perform a grid search with 7*7 different parameters with 3 different train/test splits.

This is done on a random forest classifier with 1000 trees. This will take ~30-40 mins.

```

In [19]: # turn that dataframe into a numpy array
X = Xdf.as_matrix()

# split into training and test set
for trainSplit in trainTestSplit:
    [Xtrain, Xtest, Ytrain, Ytest] = ms.train_test_split(X, Y, train_
size=trainSplit)
    scoresKey = str(trainSplit*100)[:2]
    clfKey = "Xdf_"+str(trainSplit*100)[:2]
    clf[clfKey] = gridSearchAndPlot(rand_forest,Xtrain,Ytrain)
    #(scores["precision"][scoresKey],recall["recall"][scoresKey],scores["F"][scoresKey]) = calcFscore(clf[clfKey],Xtest,Ytest)
    #print("Best parameters for trainTestSplit =",trainSplit,clf[clfKey].best_params_)

    print("\nTrain/Test Split: {:.1f}/{:.1f}".format(trainSplit, 1-trainSplit))
    print("Optimal Max Features = \t{}".format(clf[clfKey].best_params_["max_features"]))
    print("Optimal Min Samples = \t{}".format(clf[clfKey].best_params_["min_samples_split"]))

```

Train/Test Split: 0.4/0.6

Optimal Max Features = 5

Optimal Min Samples = 7

Train/Test Split: 0.6/0.4



Optimal Max Features = 3


Optimal Min Samples = 5

Train/Test Split: 0.8/0.2

Optimal Max Features = 7

Optimal Min Samples = 8





```

In [21]: # turn that dataframe into a numpy array
X = Xdf_nofe.as_matrix()

# split into training and test set
for trainSplit in trainTestSplit:
    [Xtrain, Xtest, Ytrain, Ytest] = ms.train_test_split(X, Y, train_
size=trainSplit)
    scoresKey = str(trainSplit*100)[:2]
    clfKey = "Xdf_"+str(trainSplit*100)[:2]
    clf[clfKey] = gridSearchAndPlot(rand_forest,Xtrain,Ytrain)
    #(scores["precision"][scoresKey],recall["recall"][scoresKey],scores["F"][scoresKey]) = calcFscore(clf[clfKey],Xtest,Ytest)
    #print("Best parameters for trainTestSplit =",trainSplit,clf[clfKey].best_params_)

    print("\nTrain/Test Split: {:.1f}/{:.1f}".format(trainSplit, 1-trainSplit))
    print("Optimal Max Features = \t{}".format(clf[clfKey].best_params_["max_features"]))
    print("Optimal Min Samples = \t{}".format(clf[clfKey].best_params_["min_samples_split"]))

```

Train/Test Split: 0.4/0.6

Optimal Max Features = 4

Optimal Min Samples = 7

Train/Test Split: 0.6/0.4

Optimal Max Features = 7

Optimal Min Samples = 6

Train/Test Split: 0.8/0.2

Optimal Max Features = 7

Optimal Min Samples = 8

